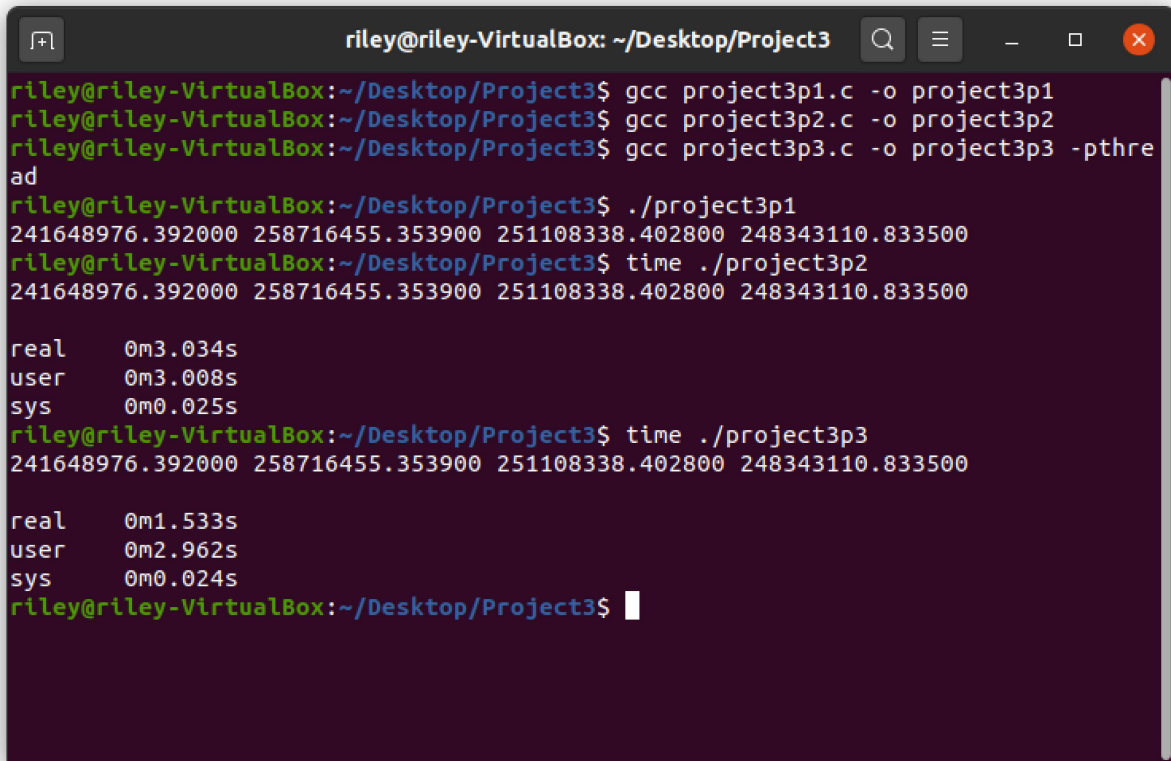


Riley Ruckman
TCES420, Au20
Project 3

Runtime of all parts (except Part1; expecting runtime to be similar to Part 2):



```
riley@riley-VirtualBox: ~/Desktop/Project3
riley@riley-VirtualBox:~/Desktop/Project3$ gcc project3p1.c -o project3p1
riley@riley-VirtualBox:~/Desktop/Project3$ gcc project3p2.c -o project3p2
riley@riley-VirtualBox:~/Desktop/Project3$ gcc project3p3.c -o project3p3 -pthread
riley@riley-VirtualBox:~/Desktop/Project3$ ./project3p1
241648976.392000 258716455.353900 251108338.402800 248343110.833500
riley@riley-VirtualBox:~/Desktop/Project3$ time ./project3p2
241648976.392000 258716455.353900 251108338.402800 248343110.833500

real    0m3.034s
user    0m3.008s
sys     0m0.025s
riley@riley-VirtualBox:~/Desktop/Project3$ time ./project3p3
241648976.392000 258716455.353900 251108338.402800 248343110.833500

real    0m1.533s
user    0m2.962s
sys     0m0.024s
riley@riley-VirtualBox:~/Desktop/Project3$
```

Part 4: The difference between parts 2 and 3 is the use of two parallel threads to perform two `matmul()` functions at the same time in part 3 instead of running them concurrently in part 2. Since part 3 has two parallel threads to do the calculations in `matmul()`, it is able to complete all `matmul()` calls in half the time of part 2, since it's using twice the number of threads as part 2. This can be seen in the “real” times for parts 2 and 3, where part 2 has 3.034s and part 3 has 1.533s.

Part 1 Code:

```
// Riley Ruckman
// TCES420, Au20
// Project 3

#include <stdio.h>
#include <stdlib.h>

double first[1000][1000], second[1000][1000], mult[1000][1000];
int m=1000, n=1000, p=1000; /* matric dimensions mxn and nxp */

void* matmul(void *);

int main(void) {

    FILE *f;

    // Read contents of matrix1.data into first
    f = fopen("matrix1.data", "r");
    fread(first, sizeof(double), sizeof(first)/sizeof(double), f);
    fclose(f);

    // Read contents of matrix2.data into second
    f = fopen("matrix2.data", "r");
    fread(second, sizeof(double), sizeof(second)/sizeof(double), f);
    fclose(f);

    matmul(NULL);

    // Read contents of matrix2.data into second
    f = fopen("matrix3.data", "w");
    fwrite(mult, sizeof(double), sizeof(mult)/sizeof(double), f);
    fclose(f);

    printf("%lf %lf %lf %lf\n", mult[6][0], mult[5][3], mult[5][4], mult[901][7]);

    return 0;
}

// Multiplies two global 1000-by-1000 matrices into 1000-by-1000 matrix
void* matmul (void *arg) {

    int i, j, k;
```

```
int s = 0, e = m;
double sum = 0;
for (i = s; i < e; i++) {
    for (j = 0; j < p; j++) {
        for (k = 0; k < n; k++) {
            sum = sum + first[i][k]*second[k][j];
        }
        mult[i][j] = sum;
        sum = 0;
    }
}
```

Part 2 Code:

```
// Riley Ruckman
// TCES420, Au20
// Project 3
```

```
#include <stdio.h>
#include <stdlib.h>
#include <pthread.h>
```

```
double first[1000][1000], second[1000][1000], mult[1000][1000];
int m=1000, n=1000, p=1000; /* matric dimensions mxn and nxp */
```

```
void* matmul(void *);
```

```
int main(void) {
```

```
    FILE *f;
```

```
    // Read contents of matrix1.data into first
    f = fopen("matrix1.data", "r");
    fread(first, sizeof(double), sizeof(first)/sizeof(double), f);
    fclose(f);
```

```
    // Read contents of matrix2.data into second
    f = fopen("matrix2.data", "r");
    fread(second, sizeof(double), sizeof(second)/sizeof(double), f);
    fclose(f);
```

```
    /*
```

```
    // Creates two threads for two calls of matmul(). Each thread is executed consecutively
to
```

```
    // simulate single-threaded performance.
    pthread_t th1, th2;
    // mult = firstLower * second + firstHigher * second
    int sec = 1;
    pthread_create(&th1, NULL, matmul, (void*)&sec);
    pthread_join(th1, NULL);
```

```
    sec = 2;
    pthread_create(&th2, NULL, matmul, (void*)&sec);
    pthread_join(th2, NULL);
    */
```

```
    int sec = 1;
```

```

    matmul((void*)&sec);
    sec = 2;
    matmul((void*)&sec);

    // Read contents of matrix2.data into second
    f = fopen("matrix3.data", "w");
    fwrite(mult, sizeof(double), sizeof(mult)/sizeof(double), f);
    fclose(f);

    printf("%lf %lf %lf %lf\n", mult[6][0], mult[5][3], mult[5][4], mult[901][7]);

    return 0;
}

// Multiplies two global 1000-by-1000 matrices into 1000-by-1000 matrix
void* matmul (void *arg) {

    int i, j, k;
    int s, e;

    if ((int *)arg == NULL) {

        s = 0;
        e = m;
    } else if (*(int *)arg == 1) {

        s = 0;
        e = m/2;
    } else if (*(int *)arg == 2) {

        s = m/2;
        e = m;
    }
    double sum = 0;
    for (i = s; i < e; i++) {
        for (j = 0; j < p; j++) {
            for (k = 0; k < n; k++) {
                sum = sum + first[i][k]*second[k][j];
            }
            mult[i][j] = sum;
            sum = 0;
        }
    }
}

```

}

Part 3 Code:

```
// Riley Ruckman
// TCES420, Au20
// Project 3

#include <stdio.h>
#include <stdlib.h>
#include <pthread.h>

double first[1000][1000], second[1000][1000], mult[1000][1000];
int m=1000, n=1000, p=1000; /* matric dimensions mxn and nxp */

void* matmul(void *);

int main(void) {

    FILE *f;

    // Read contents of matrix1.data into first
    f = fopen("matrix1.data", "r");
    fread(first, sizeof(double), sizeof(first)/sizeof(double), f);
    fclose(f);

    // Read contents of matrix2.data into second
    f = fopen("matrix2.data", "r");
    fread(second, sizeof(double), sizeof(second)/sizeof(double), f);
    fclose(f);

    // Creates two threads for two calls of matmul(). Each thread is executed in parallel for
    multi-threaded
    // performance.
    pthread_t th1, th2;

    // mult = firstLower * second + firstHigher * second
    int sec1 = 1;
    pthread_create(&th1, NULL, matmul, (void*)&sec1);

    int sec2 = 2;
    pthread_create(&th2, NULL, matmul, (void*)&sec2);

    // Waits for each thread to finish
    pthread_join(th1, NULL);
    pthread_join(th2, NULL);
}
```

```

// Read contents of matrix2.data into second
f = fopen("matrix3.data", "w");
fwrite(mult, sizeof(double), sizeof(mult)/sizeof(double), f);
fclose(f);

printf("%lf %lf %lf %lf\n", mult[6][0], mult[5][3], mult[5][4], mult[901][7]);

return 0;
}

// Multiplies two global 1000-by-1000 matrices into 1000-by-1000 matrix
void* matmul (void *arg) {

    int i, j, k;
    int s, e;

    if ((int*)arg == NULL) {

        s = 0;
        e = m;
    } else if (*(int*)arg == 1) {

        s = 0;
        e = m/2;
    } else if (*(int*)arg == 2) {

        s = m/2;
        e = m;
    }
    double sum = 0;
    for (i = s; i < e; i++) {
        for (j = 0; j < p; j++) {
            for (k = 0; k < n; k++) {
                sum = sum + first[i][k]*second[k][j];
            }
            mult[i][j] = sum;
            sum = 0;
        }
    }
    pthread_exit(NULL);
}

```